

# Yuhan Tan

(607)–233–3661 - [yt628@cornell.edu](mailto:yt628@cornell.edu) - [linkedin.com/in/yuhan-aaron-tan](https://www.linkedin.com/in/yuhan-aaron-tan)

## SKILLS

**Languages:** Java, Python, C++, C#, C, SQL, HTML, R

**Framework&Tools:** Spring, Spring Boot, MySQL, Redis, MyBatis Plus, Kafka, Docker, Linux, Knife4j, .NET Framework, Hadoop, Hive, HDFS, Qt, Git, PyTorch, TensorFlow, Numpy, OpenCV

## EDUCATION

### Cornell University

Candidate for Master of Information Science; GPA: 3.8/4.0

Ithaca, USA

08/2023 - Present

### University of Liverpool

(First Class Hons)BEng in Computer Science and Electronic Engineering; GPA: 3.8/4.0

Liverpool, UK

09/2021 - 07/2023

### Xi'an Jiaotong-Liverpool University

Major in Computer Science and Technology

Suzhou, China

09/2019 - 07/2021

## EXPERIENCE

### Software Development Intern | Spring Boot, MyBatis Plus, MySQL, Redis

05/2024 - 08/2024

NextTier

- Utilized **Redis** to implement **distributed sessions**, synchronizing login states across clusters. Used **Hash** instead of **String** to store user information, saving memory and facilitating single field modifications.
- Implemented friend similarity matching based on the **edit distance algorithm** to find the most similar users based on tags, employing a **priority queue** to **reduce memory usage** during the TOP N computation.
- Enhanced concurrency control** by implementing **Redisson distributed locks** to prevent duplicate team joining and exceeding team capacity, ensuring **mutual exclusion** and **API idempotency**.
- Optimizing caching** by storing user information lists in **Redis**, reducing API response time **from 1 second to 25 milliseconds** and ensuring data integrity with **custom Redis serializers**.
- Improved initial access speed** by implementing scheduled **cache warming** with **Spring Scheduler**.
- Enhanced database write operations** by using **custom thread pools** with **CompletableFuture** concurrency, significantly reducing import time for **1 million rows from 300 seconds to 54 seconds**.

### Software Development Intern | Kafka, Spring Boot

08/2023 - 12/2023

Eth Technology

- Developed a **streaming microservice** capable of processing over **1000 events** simultaneously using **Kafka** and **Spring Boot**.
- Self-learned** Kafka Consumer & Producer patterns in **1 month**, building **REST API** for event consumption and publication to Kafka topics.
- Implemented **unit tests** and **integration tests** using **JUnit** and **Embedded Kafka**; Conducted **end-to-end manual testing** for different scenarios of data-streaming API using **Postman**; Implemented **concurrency testing** & automated the **load tests** process using **Jmeter**.
- Integrated **Spring JPA** and utilized **H2 database** to store events metadata.

## PROJECTS

### Asynchronous Processing Framework: AaronFlow | Spring Boot, MySQL, Redis

03/2024 - Present

AaronFlow is an asynchronous task processing framework developed in Java that supports **automatic scheduling**, **automatic retries**, and **flexible task configuration**.

- Architecture Design:** Designed the application with two main layers: Flowsvr (Server) and Aaron (Worker). The Flowsvr layer provides HTTP services for **task querying**, **task scheduling**, and **task management**; the Aaron layer is responsible for **pulling and consuming tasks**.
- Database Tables Design:** Separated the storage of task information, configuration, and scheduling to reduce dependencies between tables, achieving a **loosely coupled** design. This allows for flexible **task registration** and **management**, and enables quick **task retrieval** through **indexing**.
- Task Management:** Implemented **timeout task monitoring and recovery** using a **polling** mechanism to regularly check task status, and monitored table size to trigger **table partitioning** logic when thresholds are reached (5 million records).
- Multi-Worker Optimization:** Initially used **MySQL row-level locking** to prevent multiple Workers from pulling the same batch of tasks, later introduced **Redis distributed locks** from the Worker side.
- Performance Optimization:** Conducted **stress testing** and analyzed performance bottleneck. By using a **MySQL connection pool** and increasing the maximum number of connections, improved performance **from 100 QPS to 500 QPS**.