Aaron Tan

(475)-326-9147 - aaron1004780912@163.com - linkedin.com/in/yuhan-aaron-tan

EDUCATION

University of Liverpool

Cornell University Ithaca, USA

Candidate for Master of Information Science; GPA: 3.8/4.0

08/2023 - Present Liverpool, UK

(First Class Hons) BEng in Computer Science and Electronic Engineering; GPA: 3.8/4.0

09/2021 - 07/2023

Xi'an Jiaotong-Liverpool University

Suzhou, China

Major in Computer Science and Technology

09/2019 - 07/2021

SKILLS

Languages: Java, Python, C++, C#, C, SOL, Shell, HTML, R

Framework&Tools: Spring, Spring Boot, MySQL, Redis, MyBatis Plus, Kafka, Docker, Linux, CentOS, Knife4j,

Hadoop, Hive, HDFS, VIM, Qt, GIT, PyTorch, TensorFlow, Numpy, AWS, Jira, Agile

EXPERIENCE

Software Development Intern | Spring Boot, MyBatis Plus, MySQL, Redis

05/2024 - 08/2024

NextTier

- Implemented distributed sessions to synchronize login states across distributed servers by using Redis. Enabled single field modification by using Hash instead of JSON to store user information, reducing memory by 20%.
- Reduced API response time by 95% by caching user information in Redis and ensuring data integrity with custom Redis serializers.
- Achieved scheduled cache warming with Spring Scheduler, improving initial access speed by 97.7%.
- Improved database write operations by using custom thread pools with CompletableFuture concurrency, reducing import time for 1 million rows from 300 seconds to 54 seconds.
- Introduced friend similarity matching function using the edit distance algorithm. Optimized matching speed from 34 seconds to 7 seconds through memory optimizations, selective data retrieval, and caching strategies, inspired by large-scale recommendation systems
- Enhanced concurrency control by **Redisson distributed locks** to prevent duplicate team joining and exceeding team capacity, ensuring mutual exclusion and API idempotency.

Software Development Intern | Kafka, Spring Boot

08/2023 - 12/2023

- Eth Technology
- Developed a streaming microservice capable of processing over 1000 events concurrently, using Kafka and **Spring Boot** to ensure efficient and **scalable** event handling.
- Designed and built **REST APIs** for event consumption and publication to Kafka topics, implementing Kafka Consumer and Producer patterns.
- Implemented unit tests and integration tests using JUnit and Embedded Kafka, achieving 90% code coverage; Conducted end-to-end testing for different scenarios of data-streaming APIs using Postman; Implemented concurrency testing & automatic load testing process using Jmeter.
- Integrated **Spring JPA** and utilized **H2 database** to store events metadata.

PROJECTS

Asynchronous Processing Framework: AaronFlow | Spring Boot, MySQL, Redis 03/2024 - Present AaronFlow is an asynchronous task processing framework developed in Java that supports automatic scheduling, automatic retries, and flexible task configuration.

- Implemented a two-layer architecture with Flowsvr (Server) providing HTTP services via APIs for task querying, scheduling, and management, and Aaron (Worker) pulling and consuming tasks, achieving a separation between distributed task scheduling logic and business logic.
- Designed database tables by separating task information, configuration, and scheduling, achieving a loosely coupled structure that reduces dependencies between tables. Used indexing for quick task retrieval, allowing flexible task registration and management.
- Implemented **timeout task monitoring** and **recovery**, using a **polling** mechanism to regularly check task status. Monitored table size to trigger **table partitioning** when thresholds are reached.
- Optimized multi-worker coordination by initially using MySQL row-level locking to prevent multiple Workers from pulling the same batch of tasks, and later improved performance by introducing Redis distributed locks on the Aaron (Worker) side.
- Implemented performance optimization by conducting stress testing with wrk and Lua scripts to analyze bottlenecks. Utilized a MySQL connection pool and increased the maximum number of connections, improving throughput from 100 QPS to 500 QPS.